

# Flash attention formula

## Original Attention

```
def scaled_dot_product_attention(query, key, value, attn_mask=None,
    is_causal=False, scale=None, enable_gqa=False) -> torch.Tensor:
    L, S = query.size(-2), key.size(-2)
    scale_factor = 1 / math.sqrt(query.size(-1)) if scale is None else scale
    attn_bias = torch.zeros(L, S, dtype=query.dtype)
    if is_causal:
        assert attn_mask is None
        temp_mask = torch.ones(L, S, dtype=torch.bool).tril(diagonal=0)
        attn_bias.masked_fill_(temp_mask.logical_not(), float("-inf"))
        attn_bias.to(query.dtype)

    if attn_mask is not None:
        if attn_mask.dtype == torch.bool:
            attn_bias.masked_fill_(attn_mask.logical_not(), float("-inf"))
        else:
            attn_bias += attn_mask

    key = key.repeat_interleave(query.size(-3)//key.size(-3), dim=-3)
    value = value.repeat_interleave(query.size(-3)//value.size(-3), dim=-3)

    attn_weight = query @ key.transpose(-2, -1) * scale_factor
    attn_weight += attn_bias
    attn_weight = torch.softmax(attn_weight, dim=-1)
    attn_weight = torch.dropout(attn_weight, dropout_p, train=True)
    return attn_weight @ value
```

$$S = QK^T, P = \text{softmax}(X), O = PV$$

Standard attention implementations materialize the matrices S and P to HBM, which takes  $O(N_2)$  memory.

As some or most of the operations are memory-bound (e.g., softmax), the large number of memory accesses translates to slow wall-clock time.

## Online softmax

### safe softmax

$$\frac{e^{x_i - m}}{\sum_{j=1}^N e^{x_j - m}}$$

1. for  $i \leftarrow 1, N$ 
  - $m_i = \max(m_{i-1}, x_i)$
2. for  $i \leftarrow 1, N$ 
  - $d_i = d_{i-1} + e^{x_i - m_N}$
3. for  $i \leftarrow 1, N$ 
  - $a_i = \frac{e^{x_i - m_N}}{d_N}$

### online softmax

$$\begin{aligned} d'_i &= \sum_{j=1}^i e^{x_j - m_i} = (\sum_{j=1}^{i-1} e^{x_j - m_i}) + e^{x_i - m_i} \\ &= (\sum_{j=1}^{i-1} e^{x_j - m_{i-1}}) e^{m_{i-1} - m_i} + e^{x_i - m_i} \\ &= d'_{i-1} e^{m_{i-1} - m_i} + e^{x_i - m_i} \end{aligned}$$

1. for  $i \leftarrow 1, N$ 
  - $m_i = \max(m_{i-1}, x_i)$
  - $d'_i = d'_{i-1} e^{m_{i-1} - m_i} + e^{x_i - m_i}$
2. for  $i \leftarrow 1, N$ 
  - $a_i = \frac{e^{x_i - m_N}}{d'_N}$

## Flash Attention

### Multi-pass self-attention

inputs:  $Q[k, :], K^T[:, i], V[i, :]$

output:  $O[k, :]$

1. for  $i \leftarrow 1, N$

$$x_i = Q[k, :]K^T[:, i]$$

$$m_i = \max(m_{i-1}, x_i)$$

$$d'_i = d'_{i-1}e^{m_{i-1}-m_i} + e^{x_i-m_i}$$

2. for  $i \leftarrow 1, N$

$$a_i = \frac{e^{x_i-m_N}}{d'_N}$$

$$o_i = o_{i-1} + a_i V[i, :]$$

3.  $O[k, :] = o_N$

$$o_i = \sum_{j=1}^i \left( \frac{e^{x_j-m_N}}{d'_N} V[j, :] \right)$$

$$\begin{aligned} \text{let } o'_i &= \sum_{j=1}^i \frac{e^{x_j-m_i}}{d'_i} V[j, :] \\ &= \left( \sum_{j=1}^{i-1} \frac{e^{x_j-m_i}}{d'_i} V[j, :] \right) + \frac{e^{x_i-m_i}}{d'_i} V[i, :] \\ &= \left( \sum_{j=1}^{i-1} \frac{e^{x_j-m_{i-1}}}{d'_{i-1}} \frac{e^{x_j-m_i}}{e^{x_j-m_{i-1}}} \frac{d'_{i-1}}{d'_i} V[j, :] \right) + \frac{e^{x_i-m_i}}{d'_i} V[i, :] \\ &= \left( \sum_{j=1}^{i-1} \frac{e^{x_j-m_{i-1}}}{d'_{i-1}} V[j, :] \right) \frac{d'_{i-1}}{d'_i} e^{m_{i-1}-m_i} + \frac{e^{x_i-m_i}}{d'_i} V[i, :] \\ &= o'_{i-1} \frac{d'_{i-1} e^{m_{i-1}-m_i}}{d'_i} + \frac{e^{x_i-m_i}}{d'_i} V[i, :] \end{aligned}$$

### Flash attention

1. for  $i \leftarrow 1, N$

$$x_i = Q[k, :]K^T[:, i]$$

$$m_i = \max(m_{i-1}, x_i)$$

$$d'_i = d'_{i-1}e^{m_{i-1}-m_i} + e^{x_i-m_i}$$

$$o'_i = o'_{i-1} \frac{d'_{i-1} e^{m_{i-1}-m_i}}{d'_i} + \frac{e^{x_i-m_i}}{d'_i} V[i, :]$$

2.  $O[k, :] = o'_N$

### Flash attention w/ Tiling

1. for  $i \leftarrow 1, \#tiles$

$$x_i = Q[k, :]K^T[:, (i-1)b : ib]$$


$$m_i^{(local)} = \max_{j=1}^b (x_i[j])$$

$$m_i = \max(m_{i-1}, m_i^{(local)})$$

$$d'_i = d'_{i-1} e^{m_{i-1}-m_i} + \sum_{j=1}^b e^{x_i[j]-m_i}$$

$$o'_i = o'_{i-1} \frac{d'_{i-1} e^{m_{i-1}-m_i}}{d'_i} + \sum_{j=1}^b \frac{e^{x_i[j]-m_i}}{d'_i} V[j + (i-1)b, :]$$

$$2. O[k, :] = o'_{N/b}$$



The overall SRAM memory footprint depends only on B and D and is not related to L.

---

**Algorithm 1** FLASHATTENTION

---

**Require:** Matrices  $\mathbf{Q}, \mathbf{K}, \mathbf{V} \in \mathbb{R}^{N \times d}$  in HBM, on-chip SRAM of size  $M$ .

- 1: Set block sizes  $B_c = \lceil \frac{M}{4d} \rceil, B_r = \min(\lceil \frac{M}{4d} \rceil, d)$ .
  - 2: Initialize  $\mathbf{O} = (0)_{N \times d} \in \mathbb{R}^{N \times d}, \ell = (0)_N \in \mathbb{R}^N, m = (-\infty)_N \in \mathbb{R}^N$  in HBM.
  - 3: Divide  $\mathbf{Q}$  into  $T_r = \lceil \frac{N}{B_r} \rceil$  blocks  $\mathbf{Q}_1, \dots, \mathbf{Q}_{T_r}$  of size  $B_r \times d$  each, and divide  $\mathbf{K}, \mathbf{V}$  into  $T_c = \lceil \frac{N}{B_c} \rceil$  blocks  $\mathbf{K}_1, \dots, \mathbf{K}_{T_c}$  and  $\mathbf{V}_1, \dots, \mathbf{V}_{T_c}$ , of size  $B_c \times d$  each.
  - 4: Divide  $\mathbf{O}$  into  $T_r$  blocks  $\mathbf{O}_1, \dots, \mathbf{O}_{T_r}$  of size  $B_r \times d$  each, divide  $\ell$  into  $T_r$  blocks  $\ell_1, \dots, \ell_{T_r}$  of size  $B_r$  each, divide  $m$  into  $T_r$  blocks  $m_1, \dots, m_{T_r}$  of size  $B_r$  each.
  - 5: **for**  $1 \leq j \leq T_c$  **do**
  - 6:   Load  $\mathbf{K}_j, \mathbf{V}_j$  from HBM to on-chip SRAM.
  - 7:   **for**  $1 \leq i \leq T_r$  **do**
  - 8:     Load  $\mathbf{Q}_i, \mathbf{O}_i, \ell_i, m_i$  from HBM to on-chip SRAM.
  - 9:     On chip, compute  $\mathbf{S}_{ij} = \mathbf{Q}_i \mathbf{K}_j^T \in \mathbb{R}^{B_r \times B_c}$ .
  - 10:    On chip, compute  $\tilde{m}_{ij} = \text{rowmax}(\mathbf{S}_{ij}) \in \mathbb{R}^{B_r}, \tilde{\mathbf{P}}_{ij} = \exp(\mathbf{S}_{ij} - \tilde{m}_{ij}) \in \mathbb{R}^{B_r \times B_c}$  (pointwise),  $\tilde{\ell}_{ij} = \text{rowsum}(\tilde{\mathbf{P}}_{ij}) \in \mathbb{R}^{B_r}$ .
  - 11:    On chip, compute  $m_i^{\text{new}} = \max(m_i, \tilde{m}_{ij}) \in \mathbb{R}^{B_r}, \ell_i^{\text{new}} = e^{m_i - m_i^{\text{new}}} \ell_i + e^{\tilde{m}_{ij} - m_i^{\text{new}}} \tilde{\ell}_{ij} \in \mathbb{R}^{B_r}$ .
  - 12:    Write  $\mathbf{O}_i \leftarrow \text{diag}(\ell_i^{\text{new}})^{-1} (\text{diag}(\ell_i) e^{m_i - m_i^{\text{new}}} \mathbf{O}_i + e^{\tilde{m}_{ij} - m_i^{\text{new}}} \tilde{\mathbf{P}}_{ij} \mathbf{V}_j)$  to HBM.
  - 13:    Write  $\ell_i \leftarrow \ell_i^{\text{new}}, m_i \leftarrow m_i^{\text{new}}$  to HBM.
  - 14:   **end for**
  - 15: **end for**
  - 16: Return  $\mathbf{O}$ .
-